

---

**Network Programming Laboratory**  
**Winter 2018/2019**

---

Session 1:  
Wireshark and Sockets API  
UDP Echo Server and Client

Janne Riihijärvi

# Contents

<b>Introduction</b>	<b>1</b>
<b>1 The Sockets API</b>	<b>1</b>
1.1 The Sockets Concept . . . . .	1
1.2 Address Structures . . . . .	1
1.3 Functions for Connection Control . . . . .	2
1.3.1 socket() . . . . .	3
1.3.2 close() . . . . .	3
1.3.3 bind() . . . . .	3
1.3.4 listen() . . . . .	4
1.3.5 accept() . . . . .	4
1.3.6 connect() . . . . .	4
1.4 Functions for Data Transfer . . . . .	4
1.5 Functions for Byte Ordering . . . . .	5
1.6 Functions to Handle Internet Addresses . . . . .	6
<b>2 Wireshark</b>	<b>6</b>
2.1 wireshark—New Visual Packet Sniffing . . . . .	6
<b>3 Programming Task</b>	<b>6</b>
3.1 Recommended Development Environment . . . . .	6
3.2 UDP Data Transfer . . . . .	7
3.3 Echo Server and Client . . . . .	7
<b>4 In the Lab</b>	<b>7</b>
4.1 Getting familiar with the environment . . . . .	7
4.2 Wireshark . . . . .	7
4.3 Programming . . . . .	8
<b>5 Useful Networking Tools and Commands</b>	<b>8</b>
5.1 ifconfig—Device Administration . . . . .	8
5.2 route—Route Selection . . . . .	8
5.3 ping—Testing Reachability . . . . .	8
5.4 traceroute—Packet Tracing . . . . .	9
5.5 host—Name Service . . . . .	9
5.6 netstat—Network Status . . . . .	9
5.7 tcpdump—Packet Sniffing . . . . .	9
5.8 ethereal—Visual Packet Sniffing . . . . .	9
5.9 ntop—Traffic Statistics . . . . .	11
5.10 nmap—Port Scanner . . . . .	11
<b>A Resources for Further Information</b>	<b>12</b>

## Introduction

Welcome to the first Network Programming lab session. We hope you will enjoy it and learn things you consider useful. Please spend some time studying this text carefully before coming to the lab. This preparation is helpful to solve the tasks and ensure a productive lab time.

This session will deal with three topics. In the first part we conduct some C programming exercises to brush up on the basics of the language and to become familiar with the programming environment. In the second part you will learn how to use some common networking tools under Linux, like Wireshark to capture and analyze the transmitted packages. These skills will be very useful later on in debugging programs using the network. Then you will program a client that reads strings from the keyboard and sends them over a UDP connection to another program (on your own or another computer). Then it should wait for a return message and display it on the screen. For the other side, program a server that waits for incoming messages and returns them to the sender.

## 1 The Sockets API

### 1.1 The Sockets Concept

To facilitate the programming of networked applications, the abstraction of sockets was introduced (because they were first used in the BSD version of UNIX, they are sometimes called BSD sockets). The sockets are a unified interface, over which an application can exchange data with the network. This interface looks the same for all protocols. The choice of the protocol (e.g. TCP or UDP) is specified using three parameters: *family*, *type* and *protocol*.

The Sockets API follows the UNIX paradigm to map all objects, which can be accessed to read or write, to files, so that they can be processed with the normal file I/O operations (like `read` and `write`). Sending and receiving in communications fits well to that approach. The communication endpoints are the sockets. Sending data thus means writing it to the socket (like to a file).

In the first three sessions of this lab course we will only work with TCP and UDP sockets. They both belong to the protocol family `PF_INET`, but have a different type: the type `SOCK_STREAM` specifies TCP, the type `SOCK_DGRAM` UDP.

It is also possible to directly access the IP layer (layer 3) by using raw sockets (type `SOCK_RAW`). Linux also allows an elegant method to directly access a network device (layer 2) using sockets of protocol family `PF_PACKET`.

### 1.2 Address Structures

Messages must be addressed somehow to reach a destination, and this addressing is mapped to a socket. To specify an address of any kind (generic for all protocols), the following address structure is provided:<sup>1</sup>

```
1 typedef unsigned short sa_family_t;
2 struct sockaddr {
```

---

<sup>1</sup>NOTE: All code is written in C language and is assumed to be compiled with gcc.

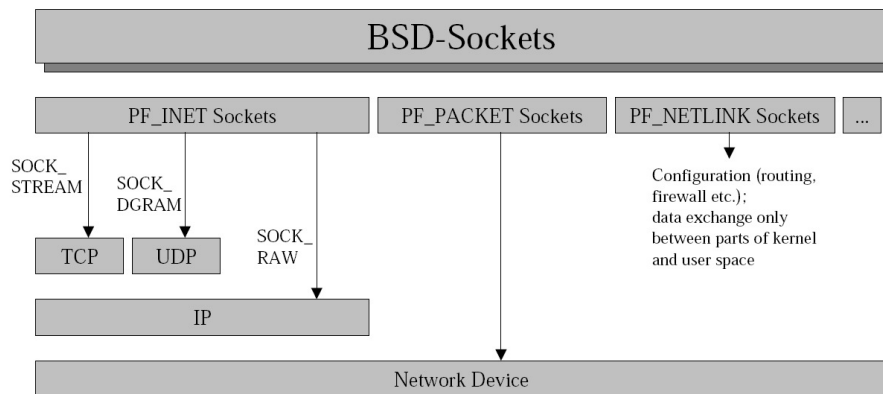


Figure 1: Structure of sockets support in Linux.

```

3  sa_family_t sa_family;
4  char        sa_data[14];
5  };

```

The element `sa_family` specifies the address family, e.g. `AF_INET` for the family of Internet protocols (`PF_INET`). The data area `sa_data[14]` is filled differently by different protocols. For IPv4 addresses the following structure `sockaddr_in` is defined:

```

1  struct in_addr { __u32 s_addr; };
2  struct sockaddr_in {
3      sa_family_t      sin_family; // Address family AF_INET
4      unsigned short int sin_port;  // Port number
5      struct in_addr    sin_addr;   // IPv4 address
6      // Pad to the size of struct socaddr
7      unsigned char     sin_zero[sizeof(struct sockaddr)
8                               - sizeof(sa_family_t)
9                               - sizeof(uint_16_t)
10                              - sizeof(struct in_addr)];
11 };

```

The protocol specific address structure is of course used for addressing, but when calling functions of the sockets API, their parameter is of type `struct sockaddr*`. Therefore, a pointer to the protocol specific address structure must be typecast to fit.

### 1.3 Functions for Connection Control

Sockets are represented at the API through socket descriptors, which can be handled like normal file descriptors (reading and writing). But because for example the establishment of a communication connection differs from opening a file, additional system calls are available for sockets.

Figure 2 shows which system calls are to be used and in which order. It distinguishes between the client role (left) and the server role (right). This

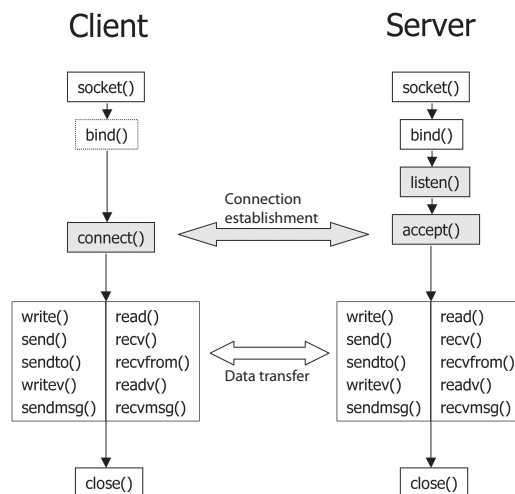


Figure 2: System calls of the sockets API. Gray shaded calls are not needed for connectionless protocols like UDP. Adapted from *Linux Netzwerkarchitektur*.

distinction does not refer to data transfer, but simply to the connection establishment: the program initiating the connection is the client, while a server is passively listening for that.

### 1.3.1 socket()

A new socket is created using the `socket()` call. The protocol (e.g. TCP or UDP) must be specified by the parameters `family`, `type` and `protocol`. Returned is a socket descriptor, that will be a parameter in future calls like `bind()`, `connect()` or `close()`. If an error occurs, `-1` is returned. The error code is then contained in `errno` and can be output for example by calling `perror()`.

```
#include <sys/types.h> // for things like AF_INET, SOCK_DGRAM
#include <sys/socket.h> // for socket()
```

```
int socket(int family, int type, int protocol);
```

### 1.3.2 close()

The normal UNIX system call `close()` is used to close a socket. Its sole parameter is the socket descriptor. `close()` returns zero on success, `-1` if an error occurred.

```
int close(int s);
```

### 1.3.3 bind()

A newly created socket doesn't have any mapping to a local address or port number. `bind()` provides such a mapping. For a client's local address it doesn't

have to be called, a port number is then selected automatically (ephemeral port number). For a server address the port number must be specified, because otherwise the client couldn't address the server.

```
int bind(int s, struct sockaddr* local_addr, int addr_len);
```

### 1.3.4 listen()

The system call `listen()` prepares a socket (on the server side) for incoming connections. The operating system will queue incoming connection requests into a socket queue. `listen()` is only used for connection oriented socket types (e.g. `SOCK_STREAM`). Parameters are the socket descriptor and the maximum number of connection requests to be queued (normally 5).

```
int listen(int s, int backlog);
```

### 1.3.5 accept()

The system call `accept()` dequeues the next connection request from the socket queue that was initiated using `listen()`. If the queue is empty when `accept()` is called, the process is blocked. It will be scheduled again by the operating system when a new connection request arrives.

After writing the source address of the connection request (client address and port number) into the address structure pointed to by `peer_addr`, the operating system creates a new socket for this connection and returns its socket descriptor as return value of `accept()`. In case of an error, `-1` is returned. The server thus has one listening socket (for incoming connection requests) and probably several connected sockets, one for each connection.

```
int accept(int s, struct sockaddr* peer_addr, int* addr_len);
```

### 1.3.6 connect()

Before a client can exchange data with a server using a connection oriented protocol (like TCP), it must create a socket (using the `socket()` call) and then establish a connection to the server using the `connect()` call. Later, when exchanging data, the server address doesn't need to be specified again.

For sockets for connectionless datagram services (e.g. UDP), the `connect()` call doesn't have to be used. For them, data can be sent directly (by specifying the server address).

```
int connect(int s, struct sockaddr* server_addr, int addr_len);
```

## 1.4 Functions for Data Transfer

### Sending Data

To send data over a socket, there are among others the following system calls (for the header files to include consult the man-pages):

- `size_t write(int s, const void* buffer, size_t len);`
- `int send(int s, const void* buffer, size_t len, int flags);`

- `int sendto(int s, const void* buffer, size_t len, int flags, const struct sockaddr* to_addr, socklen_t to_addr_len);`

The parameter `void* buffer` is the start address of a buffer that contains the data to be sent. The next parameter `size_t len` specifies the number of bytes to be sent starting from the buffer beginning. `flags` specify the transfer control—if you want to know what that can do, please refer to the man page (e.g. `man sendto`). All three functions return the number of bytes that have been sent.

While `write()` and `send()` can only be called if the socket is in the connected state, `sendto()` can always be called (because it has the destination address as parameter).

## Receiving Data

Analog to the sending functions described above, the following receive functions exist:

- `size_t read(int s, void* buffer, size_t len);`
- `int recv(int s, void* buffer, size_t len, int flags);`
- `int recvfrom(int s, void* buffer, size_t len, int flags, struct sockaddr* source_addr, socklen_t* source_addr_len);`

All three functions return the number of bytes that have been read. If there is nothing to read, the process is by default blocked by these system calls (can be changed using `flags`). It gets scheduled again by the operating system when there is something to read.

Additionally, the functions `readv()`, `writv()`, `sendmsg()` and `recvmsg()` are available for data transfer. They can be useful for example when messages consist of several parts.

## 1.5 Functions for Byte Ordering

Not all processors save the single bytes of larger types in the same order. The two ways to do it are called *little-endian* and *big-endian*. Intel processors are normally little-endian, while Sun processors are normally big-endian.

To correctly exchange data between those machines, a byte ordering has been specified. It is called the *network byte order*—in contrast to the *host byte order* (for the TCP/IP protocol family the network byte order has been defined as big-endian, but this is more of general interest). Your code should of course be able to run on all machines. Therefore several functions have been defined to map between host and network byte order. Their implementation on different machines can be different (on some machines they will do nothing), but you don't have to care about that. It is important to use the functions to make your code portable. The functions are (again consult the man pages for headers to be included):

- 32-bit host → net: `unsigned long htonl(unsigned long hostlong);`
- 16-bit host → net: `unsigned short htons(unsigned short hostshort);`

- 32-bit net  $\rightarrow$  host: `unsigned long ntohl(unsigned long netlong);`
- 16-bit net  $\rightarrow$  host: `unsigned short ntohs(unsigned short netshort);`

The functions for 32-bit integers are used for example for IPv4 addresses, while the functions for 16-bit integers are for example used for port numbers.

## 1.6 Functions to Handle Internet Addresses

Several functions are provided to handle IP addresses and DNS names. IP addresses are normally given in the dotted decimal notation as a string (like 134.130.222.125), but they are stored and processed as 32-bit unsigned integers. The following functions do the mapping:

- ASCII  $\rightarrow$  numerical:  
`int inet_aton(const char* cp, struct in_addr* inp);`
- numerical  $\rightarrow$  ASCII:  
`char* inet_ntoa(struct in_addr in);`
- unsigned long int `inet_addr(const char* cp);`

Instead of `inet_ntoa()` and `inet_aton()` the following two functions can be used. They have the advantage of supporting several address families (e.g. `AF_INET` and `AF_INET6`). Therefore it is often recommended to use them:

- `int inet_pton(int af, const char* src, void* dst);`
- `const char* inet_ntop(int af, const void* src, char* dst, size_t cnt);`

To get the IP address of your communication peer from your connected socket, the function `getpeername()` is available. With `gethostname()` you can get the DNS name (not the IP address) of your own computer. If you know the name of the communication peer, you can get its address using `gethostbyname()`.

## 2 Wireshark

### 2.1 wireshark—New Visual Packet Sniffing

Wireshark is an upgraded version of the classical tool Ethereal for capturing and analyzing network traffic. It becomes now the world's most popular network protocol analyzer. A short summary on use of Wireshark is provided through L2P. You will be using Wireshark to capture/understand the transmitted packages in a typical socket scenario.

## 3 Programming Task

### 3.1 Recommended Development Environment

We recommend using `eclipse` as the editor and the environment for debugging. The compiler is `gcc`, but `eclipse` knows how to call the compiler directly from



the user interface. If you dislike `eclipse`, you can also use other editors or IDEs like for example `gedit` and compile your code manually.

### 3.2 UDP Data Transfer

As a step towards the echo server and client, we recommend you to start with simplified programs: the client reads data from the keyboard and transmits it to the server, while the server is waiting for it and just puts it on the screen.

If you have trouble getting started, you might find the following tips useful:

#### Server

- Creates a socket and binds the local address to it (choose port number).
- For a UDP server you don't have to use `listen()` and `accept()` like for a TCP server; you can directly call `recvfrom()`.

#### Client

- Creates a socket.
- Can directly send data using `sendto()` (with the correct server address and port number).

It is normally convenient to test network applications (client and server) on your own computer. To do that you can address messages to the loopback interface (lo) with IP-address 127.0.0.1. You can then also watch the packets using for example `Wireshark` for that interface.

### 3.3 Echo Server and Client

Now extend the previous step to the echo client and server. The client should read a string from the keyboard, transmit it to the server, then wait for the reply string and put it on the screen. The server should receive the string and echo it.

## 4 In the Lab

### 4.1 Getting familiar with the environment

- In the beginning of the lab you will have some time to brush up on C basics and become familiar with Eclipse by going through basic C programming exercises. These will be circulated in the beginning of the lab session.

### 4.2 Wireshark

- Start by capturing network traffic generated by at least two different applications (you can use, for example, a web browser and the `traceroute` tools for this purpose).

- Learn how to filter the traffic displayed only to show traffic originating from your computer with a specific protocol. Try to see how `ping` and `tracert` work using this functionality.
- Use the TCP connection tracing functionality to see how a web browser establishes a connection, downloads a webpage, and closes down the connection again. Understanding this process will be very important in the coming sessions.

### 4.3 Programming

- Yes, the echo client and server, please.

## 5 Useful Networking Tools and Commands

This is a small but powerful selection of a wide variety of available tools. Understanding these commands is important for network programming applications. For the syntax please refer to the manual pages (`man`). To use most of the tools and commands, you will need root privileges.

### 5.1 `ifconfig`—Device Administration

For configuration of a network device under Linux the command `ifconfig` is used. It mainly provides activating, deactivating, and configuring of a device and its physical adapters. Protocol specific parameters like address and netmask as well as interface specific parameters like I/O port and interrupt can be set. To use a network device, it must be activated with `ifconfig` (and the adapter must be known by the kernel). The command does not store the information permanently. Upon reboot it is lost.

### 5.2 `route`—Route Selection

The `route` command is used to show and manipulate the IP routing table (static routing). Dynamic routes are set by routing protocol (e.g. Routing Information Protocol, RIP). `route` knows exactly two options, one to set static routes and one to remove them. Packets with addresses for which no directly fitting entry is found in the routing table are sent using the default entry.

### 5.3 `ping`—Testing Reachability

`ping` is used to test if a host is reachable over the network. `ping` sends an “ICMP Echo Request” packet to the specified computer and expects an “ICMP Echo Reply” back. If `ping` is complaining about an “unknown host”, then most probably the name service is not correctly configured—try pinging the IP address instead of the hostname then. If no echo is returned however, you cannot say the other computer is not reachable, because a firewall may be configured not to respond to ping. Sometimes this is done to look invisible to denial-of-service (DOS) attackers who use ping to watch a machine and launch an attack when its presence is detected.

## 5.4 **traceroute**—Packet Tracing

As the name of this program indicates, it provides two interesting pieces of information:

1. Do your packets reach their destination?
2. If not, where are they stopped?

Therefore, **traceroute** is often the next step if a computer is not reachable with **ping**. **traceroute** works by sending UDP packets with a small time to live (TTL) value towards the destination. Every host that is passed is counting the TTL down. If the TTL reaches zero, an ICMP “time exceeded” is replied. The TTL is then increased by one. Because the UDP packets are not intended to be processed, an unlikely port is chosen (can be configured). **traceroute** continues, until either ICMP “port unreachable” is returned (destination reached) or the TTL has reached its maximum (can be configured). If there is no answer within a certain duration, a “\*” is printed for that probe.

An interesting option is to use **traceroute** with the IP source routing option. With strict source routing the exact path for the packet is specified, with loose source routing a list of IP addresses can be given that must be passed.

## 5.5 **host**—Name Service

**host** is a simple utility to perform DNS lookups. It converts names to IP addresses and vice versa.

## 5.6 **netstat**—Network Status

**netstat** prints information about the Linux networking subsystem. It can print network connections, routing tables, interface statistics, masquerade connections and multicast memberships.

## 5.7 **tcpdump**—Packet Sniffing

**tcpdump** prints out the headers of packets on a network interface. The headers can be filtered by specifying a boolean expression (e.g. to only output traffic of a special protocol). **tcpdump** operates by putting the network interface card into promiscuous mode, so that every packet going across the wire is captured. Normally interface cards only capture link level frames addressed to the particular interface or to the broadcast address.

## 5.8 **ethereal**—Visual Packet Sniffing

This program provides similar functionality to **tcpdump**, but shows the information graphically. In comparison to **tcpdump** and **ntop** (see 5.9) it is normally used for detailed analysis of one datastream. It is useful for example if you want to reassemble and monitor data of a complete TCP conversation (like transporting emails)<sup>2</sup>. However, there has been no active development on **Ethereal** since the tool has been replaced by its fork “**Wireshark**”.

---

<sup>2</sup>On the command line similar information can be extracted from **tcpdump** data using the tool **tcpflow**.

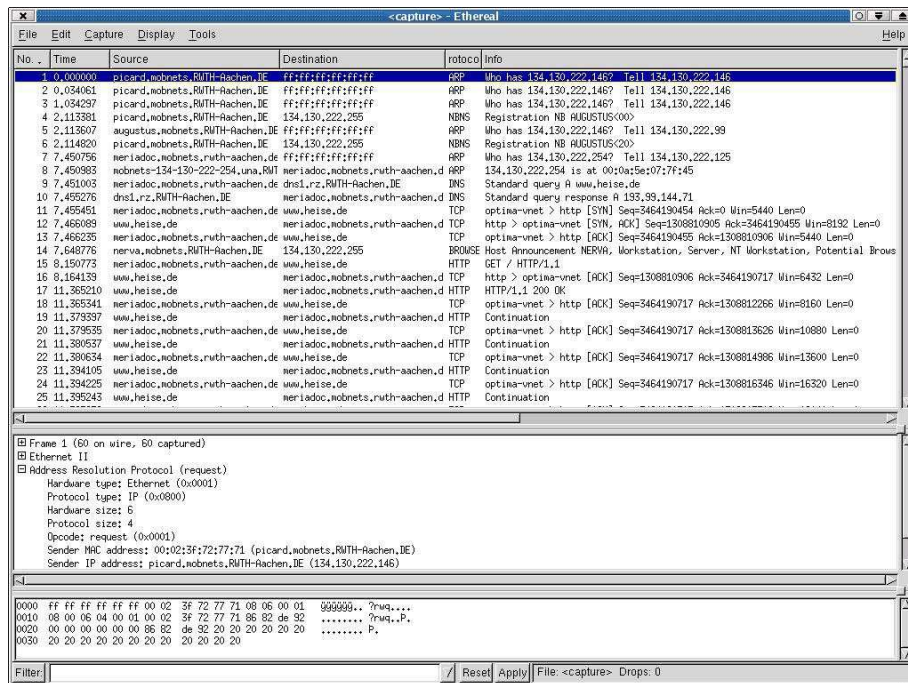


Figure 3: Packet sniffing with ethereal.

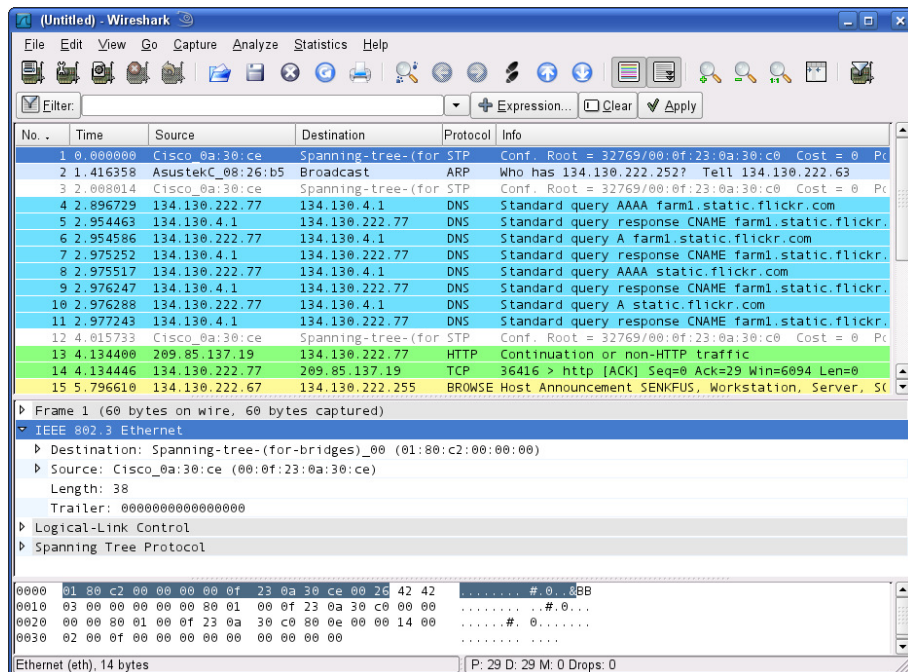


Figure 4: Packet sniffing with wireshark.

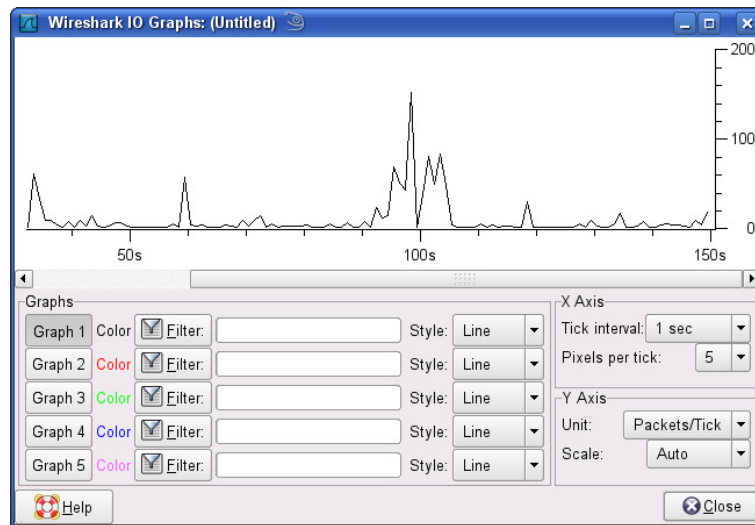


Figure 5: A throughput plot using wireshark.

## 5.9 ntop—Traffic Statistics

**ntop** is a traffic measurement and monitoring tool. It was originally programmed as a simple tool able to report the top network users (hence the name **ntop**) to quickly identify those hosts that were currently using most of the available network resources. Then it evolved into a more powerful tool. **ntop** normally works browser-based, i.e. it provides its information in the form of web pages.

example:

```
ntop -w 3000
```

In your web browser, put `http://localhost:3000`

## 5.10 nmap—Port Scanner

**nmap** is able to scan large networks to determine which hosts are up and what services they are offering. **nmap** supports a large number of scanning techniques (see the man page). It also offers a number of advanced features as for example remote OS detection via TCP/IP fingerprinting.

The result of running **nmap** is usually a list of interesting ports on the machines being scanned. **nmap** always gives the port's "well known" service name, number, state, and protocol. The state is either *open*, *filtered*, or *unfiltered*. Open means that the target machine will accept() connections on that port. Filtered means that a firewall, filter, or other network obstacle is covering the port and preventing **nmap** from determining whether the port is open. Unfiltered means that the port is known by **nmap** to be closed and no firewall/filter seems to be interfering with its attempts to determine this. Unfiltered ports are the common case and are only shown when most of the scanned ports are in the filtered state.

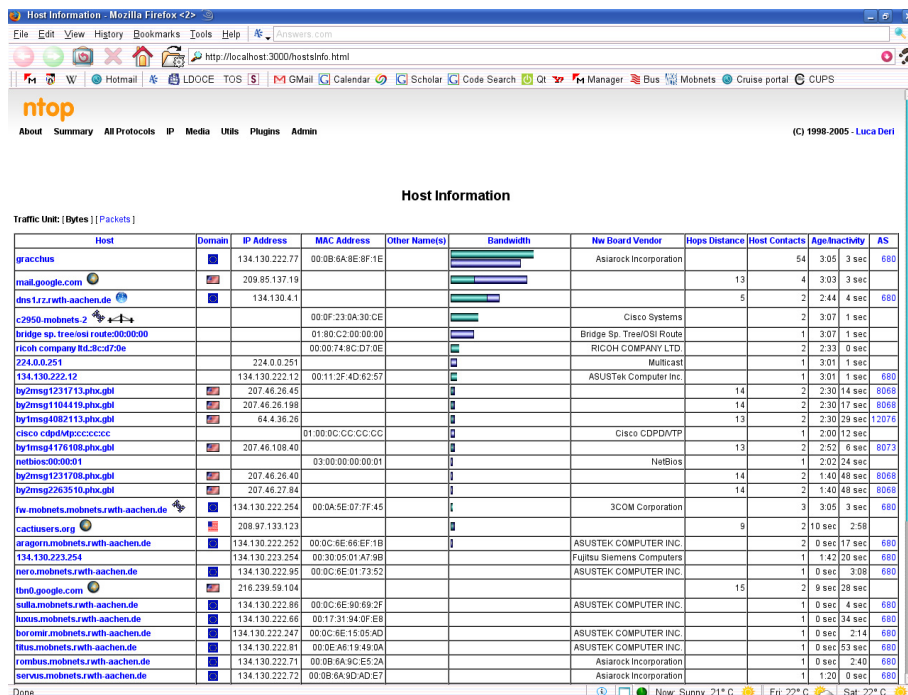


Figure 6: Traffic analysis with ntop.

## A Resources for Further Information

- Regarding the tools and commands you have to use in this session, you will most probably often rely on the man pages.
- The Internet is, of course, often a helpful resource. An interesting tutorial is for example *Beej's Guide to Network Programming: Using Internet Sockets*, which is available at <http://www.ecst.csuchico.edu/~beej/guide/net/>.

If you want to consult books, the following hints might be useful:

- The best book on network programming in C is probably *Unix Network Programming* by R. Stevens. The first of two volumes—the relevant one—is called *The Sockets Networking API*.
- Most of this session description is based on *Linux Netzwerkarchitektur* by Wehrle *et al.* The English edition covering the Linux kernel 2.6 is also available under the title *Linux Network Architecture*.
- Also very interesting and more general in scope, are *TCP/IP Illustrated Volume 1 The Protocols* and *Volume 2 The Implementation* by R. Stevens.
- For the first three sessions of this course, *Linux Socket Programming by Example* by W. Gay can be recommended.